

Filesystem Administration

Michael C. Hackett
Computer Science Department

Community
College
of Philadelphia

Lecture Topics

- Device Files
- Filesystem Types
- Hard Disk Management
 - Partitioning
 - Formatting
 - Mounting
 - Administration
- Filesystem Maintenance and Monitoring
 - Disk Usage
 - Checking for and Repairing Disk Errors
 - Mounting Filesystems at Boot
 - Disk Quotas

Device Files

- A **device file** is a file that represents a system device and specifies how to transfer data to and from the device.
- All device files are typically stored in the `/dev` directory

Device Files

- System devices are either:
 - **Character Devices**, which transfer data character-by-character to and from the device.
 - Examples: Tape drives, serial ports, parallel ports, keyboards
 - **Block Devices**, which transfer blocks of data using physical memory to buffer the transfer.
 - Examples: Hard disks and CDs/DVDs
- Block devices have faster data transfer than character devices.

Device Files

- In a long listing, character device files have a “c” before the mode.

```
[root@localhost ~]# ls -l /dev/tty1  
crw--w----. 1 gdm tty 4, 1 Dec  8 23:24 /dev/tty1
```

- Block device files have a “b” before the mode

```
[root@localhost ~]# ls -l /dev/sda1  
brw-rw----. 1 root disk 8, 1 Dec  8 23:24 /dev/sda1
```

Device Files

- Common device files:

- `/dev/fd0` First floppy disk drive
- `/dev/sr0` First CD/DVD device
 - There is often a `/dev/cdrom` file linked to `/dev/sr0`
- `/dev/hda1` First partition on the first PATA hard disk drive
- `/dev/hdb1` First partition on the second PATA hard disk drive
- `/dev/hdc1` First partition on the third PATA hard disk drive
- `/dev/hdd1` First partition on the fourth PATA hard disk drive

Device Files

- Common device files:

- `/dev/sda1` First partition on the first SCSI/SATA hard disk drive
- `/dev/sdb1` First partition on the second SCSI/SATA hard disk drive
- `/dev/tty1` First local terminal
- `/dev/tty2` Second local terminal
- `/dev/ttyS0` First serial port (COM1)
- `/dev/ttyS1` Second serial port (COM2)
- `/dev/lp0` First parallel port (LPT1)
- `/dev/st0` First SCSI tape drive

Device Files

- When performing a long listing of a device file, the file size is replaced with two numbers.
- These indicate the major and minor numbers.
 - **Major number:** Indicates the device driver for it in the kernel
 - **Minor number:** Indicates the device itself

```
[root@localhost ~]# ls -l /dev/sda1  
brw-rw----. 1 root disk 8, 1 Dec  8 23:24 /dev/sda1
```


Device Files

- In this example, the system's first SATA hard disk drive (sda) has two partitions (sda1 and sda2)

```
[root@localhost ~]# ls -l /dev/sda*  
brw-rw----. 1 root disk 8, 0 Dec  8 23:24 /dev/sda  
brw-rw----. 1 root disk 8, 1 Dec  8 23:24 /dev/sda1  
brw-rw----. 1 root disk 8, 2 Dec  8 23:24 /dev/sda2  
[root@localhost ~]#
```

- All three use the same kernel device driver, indicated by 8 (the major number)
- Each is identified by a unique minor number (0, 1, and 2)

Device Files

- Multiple devices can share the same major number.
 - Though, each device will have a unique minor number.
- If this system had two SATA hard disks, each with two partitions, the major and minor numbers for their device files might be something like:
 - sda (8, 0)
 - sda1 (8, 1)
 - sda2 (8, 2)
 - sdb (8, 3)
 - sdb1 (8, 4)
 - sdb2 (8, 5)

Device Files

- If a device file ever becomes corrupted, it will be listed as a regular file.
 - Can be identified by using **find /dev -type f**
- To recreate a device file (if it becomes corrupted or is accidentally deleted), the **mknod** command will **make** a new device file and with a corresponding **inode**

Device Files

- The syntax of the mknod command is
mknod *file type major minor*
- If tty4 were accidentally deleted it could be recreated with the following command: **mknod /dev/tty4 c 4 4**
 - Note: You need to know the device's major and minor number.

```
[root@localhost ~]# ls -l /dev/tty4
crw--w----. 1 root tty 4, 4 Dec  8 23:24 /dev/tty4
[root@localhost ~]# rm /dev/tty4
rm: remove character special file '/dev/tty4'? y
[root@localhost ~]# ls -l /dev/tty4
ls: cannot access '/dev/tty4': No such file or directory
[root@localhost ~]# mknod /dev/tty4 c 4 4
[root@localhost ~]# ls -l /dev/tty4
crw-r--r--. 1 root root 4, 4 Dec  9 00:11 /dev/tty4
[root@localhost ~]#
```

Device Files

- A listing of all devices (character and block devices) and their major number is in the `/proc/devices` file.
 - Can be viewed using: **cat /proc/devices**
- A listing of all block devices is in the `/sys/block` directory.
 - Can be viewed using: **ls /sys/block**
- To list detailed information about all **block** devices, use the **lsblk** command.

```
[root@localhost ~]# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
sda          8:0    0   30G  0 disk
├─sda1       8:1    0   3.7G  0 part [SWAP]
└─sda2       8:2    0  26.3G  0 part /
sr0         11:0    1 1024M  0 rom
```

[root@localhost ~]# _

Filesystem Types

- There are a number of different filesystem types available to use in Linux.
- Each has their own strengths and weaknesses.
- Luckily, a Linux system can use several devices formatted with different filesystems.

Filesystem Types

- FAT – File Allocation Table
 - Older filesystem compatible with many different operating systems.
 - Used for compatibility reasons; Does not provide the same capabilities as more modern filesystems.
- VFAT – Virtual FAT
 - FAT with long filename support

Filesystem Types

- exFAT – Extended FAT
 - Improved version of FAT with large file support.
 - Most commonly used filesystem for USB flash drives
- NTFS – New Technology File System
 - Proprietary Microsoft filesystem
 - Not natively supported by Linux

Filesystem Types

- ext2 (Second Extended Filesystem)
 - Improved version of the original extended filesystem
 - Traditional filesystem used in Linux
 - Still supported by Linux, but was mostly replaced by ext3
- ext3 (Third Extended Filesystem)
 - Improvement on ext2 and allows for journaling.
 - Faster startup and recovery time.
 - Insures data integrity better than ext2

Filesystem Types

- ext4 (Fourth Extended Filesystem)
 - Improved version of ext3
 - Larger filesystem support and speed enhancements
 - Supports volumes up to 1 exbibyte and files up to 16 tebibyte in size.
- xfs (X File System)
 - High-performance file system
 - Used when there is need to quickly write large numbers of files to the disk.

Filesystem Types

- btrfs (B-tree File System) (*“Butter FS”*)
 - Relatively new; intended to replace ext4
 - Geared toward large-scale storage
 - Has the ability to span multiple devices

Filesystem Types

- Those are a small selection of mostly general-purpose filesystems.
- Some filesystems function as network protocols to share data across a network.
 - Network File Systems


Filesystem Types

- **SMB** (Server Message Block)
 - Primarily used with Windows computers
 - SMB-compatible software called Samba helps interface Linux and Windows network shares
- **CIFS** (Common Internet File System)
 - Rarely used
 - Was intended to replace Samba 1, but Samba versions 2 and 3 superseded it.
- **NFS** (Network File System)
 - Preferred for network file serving between Linux clients and Linux servers
 - SMB is better for Linux/Windows network file serving

Partitioning

- Physically, hard disks contain one or more circular metal *platters* that spin at high speeds.
- Data is read from concentric circles on the platter called **tracks**
- Each track is divided into **sectors** of data
 - Sectors are combined into one or more **blocks**

Partitioning




Partitioning

- Hard disks are often split up into multiple smaller drives called **partitions**.
- Each partition contains its own filesystem and can each be mounted to different mount point directories.
 - More on this later in the lecture
- Segregating data into separate areas of the hard disk is useful from an organizational standpoint.

Partitioning

- Hard disks with a Master Boot Record (MBR) normally contain up to four **primary partitions**
 - Any of the four primary partitions can be replaced by an **extended partition**
 - Extended partitions can contain an unlimited number of (but normally up to 12) partitions called **logical partitions**
- Hard disks with a GUID Partition Table (GPT) can have up to 128 primary partitions
 - No need for extended partitions and logical drives

Partitioning



Partitioning

- Some frequently used command line utilities for partitioning hard disks:
 - **fdisk** – Fixed Disk
 - **cfdisk** – Curses **fdisk**
 - Curses is a terminal control library giving it a different interface than fdisk
 - **parted** – GNU **P**artition **E**ditor
- We'll focus on using **fdisk**

Partitioning

- The **fdisk** utility is a command line, menu-driven program that allows modifying, creating, and deleting partitions on a storage drive.

fdisk [options] *device_file*

- For example, to start the utility to edit/create/delete partitions on the second SATA hard drive:

fdisk /dev/sdb

Partitioning

```
[root@localhost ~]# fdisk /dev/sdb

Welcome to fdisk (util-linux 2.34).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.

Device does not contain a recognized partition table.
Created a new DOS disklabel with disk identifier 0x1d50330c.

Command (m for help): _
```

This was a brand new virtual disk attached as a second hard drive in the virtual machine

Partitioning

- Common **fdisk** options:
 - **-b** Specifies the number of drive sectors
 - **-H** Specifies the number of drive heads
 - **-S** Specifies the number of sectors per track
 - **-s** Print the partition size in blocks
 - **-l** List the partition table for the device
- We will use the program's menus instead of using the options.

Partitioning

- Common **fdisk** menu options:
 - **n** Create a **n**ew partition
 - **d** **D**delete a partition
 - **p** List the existing **p**artitions
 - **w** **W**rite the changes to the drive and exit
 - **m** Help **m**anual
 - **q** **Q**uit/exit fdisk without writing changes to the drive
- This is what we will use in the **fdisk** program

Partitioning

- Using the **p** menu option on a new disk lists no partitions, as none were created yet

```
[root@localhost ~]# fdisk /dev/sdb

Welcome to fdisk (util-linux 2.34).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.

Device does not contain a recognized partition table.
Created a new DOS disklabel with disk identifier 0x1d50330c.

Command (m for help): p
Disk /dev/sdb: 20.13 GiB, 21613379584 bytes, 42213632 sectors
Disk model: VBOX HARDDISK
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x1d50330c

Command (m for help): _
```


Partitioning

- Using the **n** menu option will begin creating a new partition.
 - We can then choose to make this a primary or extended partition

```
Command (m for help): n
Partition type
   p   primary (0 primary, 0 extended, 4 free)
   e   extended (container for logical partitions)
Select (default p): _
```

Partitioning

- After choosing to make it a primary partition, the option is given to identify the partition number. **n** menu option will begin creating a new partition.
 - Can be given a number 1-4 (All four are available since no others were created yet)

```
Command (m for help): n
Partition type
   p   primary (0 primary, 0 extended, 4 free)
   e   extended (container for logical partitions)
Select (default p): p
Partition number (1-4, default 1): _
```

Partitioning

- Then, you choose the first disk sector of the partition.
 - The default will be the first available sector
- Next, you choose the last sector of the partition.
 - By default, it suggests the last available sector which would take up the rest of the disk.
 - Alternatively you can specify the size in kilobytes, megabytes, gigabytes, etc.
- For this example, I will start at the first available sector up through 5 gigabytes worth of sectors.

Partitioning

```
Command (m for help): n
Partition type
  p   primary (0 primary, 0 extended, 4 free)
  e   extended (container for logical partitions)
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-42213631, default 2048): 2048
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2048-42213631, default 42213631): +5G

Created a new partition 1 of type 'Linux' and of size 5 GiB.

Command (m for help): _
```

Partitioning

- Using the **p** menu option will display this new partition

```
Command (m for help): p
Disk /dev/sdb: 20.13 GiB, 21613379584 bytes, 42213632 sectors
Disk model: UBOX HARDDISK
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x1d50330c
```

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdb1		2048	10487807	10485760	5G	83	Linux

```
Command (m for help):
```

Partitioning

- To write this change to the hard disk, the **w** menu option is used
 - **fdisk** is exited when the partition table is written

```
Command (m for help): w
The partition table has been altered.
Calling ioctl() to re-read partition table.
Syncing disks.

[root@localhost ~]# _
```

Formatting

- Before data can be stored to a disk/partition, it must be formatted with a file system.
- The **mkfs** command is used to **make** a **filesystem** on a partition.
 - The most common option used is the **-t** option, which specifies the type of filesystem.

Formatting

```
[root@localhost ~]# mkfs -t ext2 /dev/sdb1
mke2fs 1.45.3 (14-Jul-2019)
Creating filesystem with 1310720 4k blocks and 327680 inodes
Filesystem UUID: 834fe96b-8686-4cef-961a-bf1744d57c1b
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736

Allocating group tables: done
Writing inode tables: done
Writing superblocks and filesystem accounting information: done

[root@localhost ~]# _
```


Formatting

- An alternative is to use **mkfs.*filesystemtype***

```
[root@localhost ~]# mkfs.ext4 /dev/sdb1
mke2fs 1.45.3 (14-Jul-2019)
/dev/sdb1 contains a ext2 file system
    created on Sat Jan 11 01:17:23 2020
Proceed anyway? (y,N) y
Creating filesystem with 1310720 4k blocks and 327680 inodes
Filesystem UUID: e6244eb0-87db-409d-8663-dd612c68e246
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done

[root@localhost ~]#
```

Formatting

- Formatting a partition will cause any existing data on the partition to become inaccessible.
 - Specialized recovery software would be needed to retrieve the lost data

```
[root@localhost ~]# mkfs.ext4 /dev/sdb1
mke2fs 1.45.3 (14-Jul-2019)
/dev/sdb1 contains a ext2 file system
    created on Sat Jan 11 01:17:23 2020
Proceed anyway? (y,N) y
Creating filesystem with 1310720 4k blocks and 327680 inodes
Filesystem UUID: e6244eb0-87db-409d-8663-dd612c68e246
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done
```

Mounting

- Before a partition can be used in a Linux system, the partition must be formatted with a filesystem and then must be ***mounted***
 - The term originated when data was stored on large reels of magnetic tape that needed to be physically lifted/attached onto computers.
- Today, it refers to when a storage device is made accessible to the system.

Mounting

- Where the storage device is mounted is referred to its **mount point**
 - A directory in the directory tree
- When users create files and subdirectories in the mount point, the files and subdirectories are created *in the filesystem of the mounted device*.

Mounting



Mounting

- When the system is first booted, a filesystem must be mounted to the / directory
 - We know this as the root directory, but the filesystem mounted there is the **root filesystem**
- Other devices can be mounted to other directories at boot time
 - We'll come back to this soon

Mounting

- The **mount** command is used to mount a storage device
mount *devicefile mountpoint*
- The command **mount /dev/sdb1 /mnt** would mount the sdb1 partition to the /mnt directory

```
[root@localhost ~]# mount /dev/sdb1 /mnt  
[root@localhost ~]# _
```

Mounting

- In this example, any data stored in /mnt will be written to sdb1

```
[root@localhost mnt]# touch testfile  
[root@localhost mnt]# ls  
lost+found testfile  
[root@localhost mnt]#
```


Mounting

- To **unmount** a device, use the **umount** command
umount *mountpoint*
- A drive in use will not be able to be unmounted
 - This includes if you are currently in (or in a subdirectory of) the mount point

Mounting

```
[root@localhost mnt]# ls
lost+found  testfile
[root@localhost mnt]# umount /mnt
umount: /mnt: target is busy.
[root@localhost mnt]# cd
[root@localhost ~]# umount /mnt
[root@localhost ~]# ls /mnt
[root@localhost ~]# _
```

Mounting

- If I were to remount the disk, the files stored on the disk will be accessible again.

```
[root@localhost ~]# umount /mnt
[root@localhost ~]# ls /mnt
[root@localhost ~]# mount /dev/sdb1 /mnt
[root@localhost ~]# ls /mnt
lost+found  testfile
[root@localhost ~]#
```

Mounting

- Some important notes:

- Disks that are currently mounted cannot be formatted. The disk must be unmounted when it is to be formatted.
- Disks that are currently mounted cannot have their partition table changed. The disk must be unmounted when its partitions are to be added/modified/deleted.

Administration (File System Info)

- The **/proc/mounts** file contains a list of active mount points

```
[root@localhost ~]# cat /proc/mounts
sysfs /sys sysfs rw,seclabel,nosuid,nodev,noexec,relatime 0 0
proc /proc proc rw,nosuid,nodev,noexec,relatime 0 0
devtmpfs /dev devtmpfs rw,seclabel,nosuid,size=999340k,nr_inodes=249835,mode=755 0 0
securityfs /sys/kernel/security securityfs rw,nosuid,nodev,noexec,relatime 0 0
tmpfs /dev/shm tmpfs rw,seclabel,nosuid,nodev 0 0
devpts /dev/pts devpts rw,seclabel,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=000 0 0
tmpfs /run tmpfs rw,seclabel,nosuid,nodev,mode=755 0 0
cgroup2 /sys/fs/cgroup cgroup2 rw,seclabel,nosuid,nodev,noexec,relatime,nsdelegate 0 0
pstore /sys/fs/pstore pstore rw,seclabel,nosuid,nodev,noexec,relatime 0 0
bpf /sys/fs/bpf bpf rw,nosuid,nodev,noexec,relatime,mode=700 0 0
configfs /sys/kernel/config configfs rw,nosuid,nodev,noexec,relatime 0 0
/dev/sda2 / xfs rw,seclabel,relatime,attr2,inode64,logbufs=8,logbsize=32k,noquota 0 0
selinuxfs /sys/fs/selinux selinuxfs rw,relatime 0 0
systemd-1 /proc/sys/fs/binfmt_misc autofs rw,relatime,fd=29,pgrp=1,timeout=0,minproto=5,maxproto=5,direct,pipe_ino=19479 0 0
mqueue /dev/mqueue mqueue rw,seclabel,nosuid,nodev,noexec,relatime 0 0
debugfs /sys/kernel/debug debugfs rw,seclabel,nosuid,nodev,noexec,relatime 0 0
hugetlbfs /dev/hugepages hugetlbfs rw,seclabel,relatime,pagesize=2M 0 0
fusectl /sys/fs/fuse/connections fusectl rw,nosuid,nodev,noexec,relatime 0 0
tmpfs /tmp tmpfs rw,seclabel,nosuid,nodev 0 0
sunrpc /var/lib/nfs/rpc_pipefs rpc_pipefs rw,relatime 0 0
tmpfs /run/user/42 tmpfs rw,seclabel,nosuid,nodev,relatime,size=203520k,mode=700,uid=42,gid=42 0 0
tmpfs /run/user/0 tmpfs rw,seclabel,nosuid,nodev,relatime,size=203520k,mode=700 0 0
/dev/sdb1 /mnt ext4 rw,seclabel,relatime 0 0
[root@localhost ~]# _
```

Administration (File System Info)

- The **/proc/partitions** file contains a list of active partitions

```
[root@localhost ~]# cat /proc/partitions
major minor  #blocks  name

   8         0   31457280 sda
   8         1    3905536 sda1
   8         2   27550720 sda2
   8        16   21106816 sdb
   8        17    5242880 sdb1
  11         0    1048575 sr0
[root@localhost ~]#
```

Administration (File System Info)

- The **lsblk** command lists all **block** storage devices

```
[root@localhost ~]# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
sda          8:0    0   30G  0 disk
├─sda1       8:1    0   3.7G  0 part [SWAP]
└─sda2       8:2    0  26.3G  0 part /
sdb          8:16   0  20.1G  0 disk
└─sdb1       8:17   0    5G   0 part /mnt
sr0         11:0    1 1024M  0 rom
[root@localhost ~]# _
```

Administration (File System Info)

- The **-f** option displays additional information about the block device

```
[root@localhost ~]# lsblk -f
```

NAME	FSTYPE	LABEL	UUID	FSAVAIL	FSUSE%	MOUNTPOINT
sda						
└─sda1	swap		84ca921a-4a87-4cec-a24b-0f896b2f5a92			[SWAP]
└─sda2	xf		794f40cf-485b-4218-bc7a-5759e39bce5f	19.8G	25%	/
sdb						
└─sdb1	ext4		e6244eb0-87db-409d-8663-dd612c68e246	4.6G	0%	/mnt
sr0						

```
[root@localhost ~]# _
```


Administration (File System Info)

- The **dumpe2fs** command displays information about an ext 2/3/4 file system.

dumpe2fs [options] *device_file*

- Commonly used options:
- **-x** Prints a report about block numbers in the file system
- **-b** Prints the bad blocks in the file system

Administration (File System Info)

- The **xfs_info** command displays information about an xfs file system.

xfs_info [options] *device_file*

Administration (Resize File System)

- The **resize2fs** command allows for resizing an ext 2/3/4 file system.

resize2fs *device_file size*

resize2fs /dev/sdb1 4G

- An ext file system can be enlarged without unmounting it
 - Must be unmounted to shrink the file system
- Does **not** resize partitions, only the file system.

Administration (Resize File System)

- The **xfs_growfs** command allows for resizing an xfs file system.

xfs_growfs *device_file*

xfs_growfs /dev/sdb1

- To extend to a specified size:

xfs_growfs *device_file* **-D** *number_of_blocks*

Disk Usage

- The two commands commonly used for displaying disk space utilization are **df** and **du**
- The **df** command displays the **d**isk's **f**ree space
- The **du** command displays the **d**isk's **u**sage

Disk Usage

- The **df** command (without any options)

```
[root@localhost ~]# df
Filesystem      1K-blocks    Used Available Use% Mounted on
devtmpfs        999340         0    999340   0% /dev
tmpfs           1017600         0    1017600   0% /dev/shm
tmpfs           1017600     1144    1016456   1% /run
/dev/sda2       27537268 6845980    20691288 25% /
tmpfs           1017600        76    1017524   1% /tmp
tmpfs           203520      5800     197720   3% /run/user/42
tmpfs           203520         0     203520   0% /run/user/0
/dev/sdb1       5095040     20472     4796040   1% /mnt
[root@localhost ~]# _
```

Disk Usage

- The **df** command with the **-H** (“human readable”) option displays the space using byte units

```
[root@localhost ~]# df -H
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        1.1G   0    1.1G   0% /dev
tmpfs           1.1G   0    1.1G   0% /dev/shm
tmpfs           1.1G  1.2M   1.1G   1% /run
/dev/sda2       29G   7.1G   22G   25% /
tmpfs           1.1G   78k   1.1G   1% /tmp
tmpfs          209M   6.0M   203M   3% /run/user/42
tmpfs          209M    0    209M   0% /run/user/0
/dev/sdb1       5.3G   21M   5.0G   1% /mnt
[root@localhost ~]#
```

Disk Usage

- **df** is good for an overall picture of the drives/partitions
- For more granular disk utilization details, we use the **du** command.

du [**options**] *directory*

- The **du** command recursively traverses the directory
 - If a directory is not specified, the current location is traversed.

Disk Usage

- Utilization of the /boot directory:

```
[root@localhost ~]# du /boot
7080    /boot/grub2/themes/system
7080    /boot/grub2/themes
3068    /boot/grub2/i386-pc
2504    /boot/grub2/fonts
12668   /boot/grub2
0       /boot/efi/EFI/fedora
0       /boot/efi/EFI
0       /boot/efi
8       /boot/loader/entries
8       /boot/loader
141660  /boot
[root@localhost ~]#
```

Disk Usage

- Use the **-h** option for human readable sizes:

```
[root@localhost ~]# du -h /boot
7.0M    /boot/grub2/themes/system
7.0M    /boot/grub2/themes
3.0M    /boot/grub2/i386-pc
2.5M    /boot/grub2/fonts
13M     /boot/grub2
0       /boot/efi/EFI/fedora
0       /boot/efi/EFI
0       /boot/efi
8.0K    /boot/loader/entries
8.0K    /boot/loader
139M    /boot
```

Disk Usage

- Use the **-s** option for a summary of the directory:

```
[root@localhost ~]# du -s /boot
141660  /boot
[root@localhost ~]# du -hs /boot
139M    /boot
[root@localhost ~]# _
```

Disk Checking and Repair

- File system errors are commonly caused by power failures, hardware failures, and the improper shutdown of the system.
- Disk checks can be performed at start-up (every time or every *X* number of days) or performed manually.
- Disk checks are important to maintaining the integrity (correctness and validity) of the file system.

Disk Checking and Repair

- The **fsck** command is used to perform a **file system check** and can be used to correct any errors.

fsck [**options**] *device_file*

- The drive must be unmounted before the check can be performed.

```
[root@localhost ~]# fsck /dev/sdb1
fsck from util-linux 2.34
e2fsck 1.45.3 (14-Jul-2019)
/dev/sdb1 is mounted.
e2fsck: Cannot continue, aborting.
```

Disk Checking and Repair

```
[root@localhost ~]# umount /mnt
[root@localhost ~]# fsck /dev/sdb1
fsck from util-linux 2.34
e2fsck 1.45.3 (14-Jul-2019)
/dev/sdb1: clean, 12/327680 files, 42078/1310720 blocks
[root@localhost ~]#
```

Disk Checking and Repair

- Use the **-f** option to perform a full check:

```
[root@localhost ~]# fsck -f /dev/sdb1
fsck from util-linux 2.34
e2fsck 1.45.3 (14-Jul-2019)
Pass 1: Checking inodes, blocks, and sizes
Pass 2: Checking directory structure
Pass 3: Checking directory connectivity
Pass 4: Checking reference counts
Pass 5: Checking group summary information
/dev/sdb1: 12/327680 files (0.0% non-contiguous), 42078/1310720 blocks
[root@localhost ~]# _
```

Disk Checking and Repair

- If **fsck** finds any problems, it will ask you if it can fix the error.
- To avoid these prompts and allow **fsck** to automatically repair any errors, use the **-y** option.
- If there are files it cannot repair, they will be placed in the **lost+found** directory of the file system.
 - Only ext filesystems have a lost+found directory

Disk Checking and Repair

- The **fsck** command automatically calls the appropriate tool to check that file system type.
- For ext filesystems, it uses **e2fsck**
- For xfs filesystems, it uses **xfs_repair**

Disk Checking and Repair

- File systems have numerous parameters that can be configured or “tuned”
- For ext filesystems, the **tune2fs** command is used.
- For xfs filesystems, the **xfs_admin** command is used.

Disk Checking and Repair

- **tune2fs** has numerous options, but we'll focus on two:
- **-i** Specifies the number of days/months/weeks between automatic filesystem checks
- **-c** Specifies the maximum number of times the drive can be mounted between automatic file system checks

Disk Checking and Repair

- **tune2fs -i 2w /dev/sdb1** will force an automatic filesystem check on /dev/sdb1 every two weeks.

```
[root@localhost ~]# tune2fs -i 2w /dev/sdb1
tune2fs 1.45.3 (14-Jul-2019)
Setting interval between checks to 1209600 seconds
```

Disk Checking and Repair

- **tune2fs -c 5 /dev/sdb1** will force an automatic filesystem check on /dev/sdb1 after every 5 mounts.

```
[root@localhost ~]# tune2fs -c 5 /dev/sdb1
tune2fs 1.45.3 (14-Jul-2019)
Setting maximal mount count to 5
[root@localhost ~]#
```

Mounting Filesystems at Boot

- The `/etc/fstab` file contains the **filesystem table** that is used to mount devices at boot time.
- Each entry has six fields:
 - The device to mount
 - The device's mount point
 - The filesystem type of the device
 - Mounting options
 - Dump number
 - Fsck number

Mounting Filesystems at Boot

- Device to mount:
 - Devices can be identified by one of the following:
 - The path to a device file
 - The filesystem's UUID
 - The GPT partition UUID (PARTUUID)
 - The filesystem's label

Mounting Filesystems at Boot

- Device's mount point:
 - Simply specifies where to mount the filesystem
- The filesystem type of the device
 - For example, ext4 or xfs

Mounting Filesystems at Boot

- Mounting options
 - Options that would be provided if using the mount command manually (See the manual page for the mount command for the full list)
 - Commonly used options:
 - **defaults** Mounts the filesystem using default settings
 - **ro** Mounts the filesystem in a read only state
 - **noauto** Prevents the filesystem from being mounted at boot
 - Multiple options are separated by commas

Mounting Filesystems at Boot

- Dump number (0 or 1):
 - 0 indicates the filesystem should not be backed up
 - 1 indicates the filesystem should be backed up
- Fck number (0, 1, or 2):
 - 0 indicates the filesystem should not be checked
 - 1 indicates the filesystem to be checked first
 - Usually the root filesystem
 - 2 indicates the filesystem(s) to be checked next
 - Checked in the order they appear in the fstab file

Mounting Filesystems at Boot

- The `/etc/fstab` file is also checked when issuing the `mount` command from the terminal.
- The `mount` command will use the options specified for that device (if it exists) in the `/etc/fstab` file.

Mounting Filesystems at Boot

- The `/etc/fstab` file can be edited using a text editor like `vi`.
 - Must have administrator privileges to save any changes.
- A command like **`cat`** can be used to display the contents of `/etc/fstab`

Mounting Filesystems at Boot

```
[root@localhost ~]# cat /etc/fstab
#
# /etc/fstab
# Created by anaconda on Sat Nov 30 13:38:45 2019
#
# Accessible filesystems, by reference, are maintained under '/dev/disk/'.
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info.
#
# After editing this file, run 'systemctl daemon-reload' to update systemd
# units generated from this file.
#
UUID=794f40cf-485b-4218-bc7a-5759e39bce5f /
UUID=84ca9211-4a87-4cec-a24b-0f896b2f5a92 none
[root@localhost ~]#
```

Device

Mount Point

FS Type

Options

Dump #

Fsck #

Mounting Filesystems at Boot

- Adding the following line to /etc/fstab

/dev/sdb1 /mnt ext4 defaults 0 0

- Would mount /dev/sdb1 to /mnt at boot time.
 - Filesystem specified as ext4
 - No additional options set
 - No dump or checking set

Disk Quotas

- Quotas allow an administrator to allot and monitor the file system space that a user may use.
- The tools most commonly used are:

quotacheck	Creates the quota database files for a file system
edquota	Edits quotas
setquota	Sets quotas
repquota	Displays a quota report
quota	Allows regular users to check their quotas

Disk Quotas

- Quotas can restrict:
 - How many files and directories a user may create on a filesystem
 - The total size of all files owned by the user on a filesystem
- A **soft limit** is a quota that the user may exceed for some period of time (the default is seven days)
- A **hard limit** is a quota that the user may not exceed.